

SUBQUERIES

FLASHBACK: SO FAR WE HAVE SEEN
QUERIES AS STANDALONE COMMANDS
THAT FETCH DATA FROM A DATABASE

BUT IN REALITY, QUERIES ARE PRETTY
PLUG-AND-PLAY

BUT IN REALITY, QUERIES ARE PRETTY
PLUG-AND-PLAY

A QUERY IS A COMMAND THAT
RETURNS A TABLE (ROWS AND
COLUMNS)

BUT IN REALITY, QUERIES ARE PRETTY
PLUG-AND-PLAY

WE COULD CALCULATE THE UNION OF 2
QUERIES

BUT IN REALITY, QUERIES ARE PRETTY
PLUG-AND-PLAY
WE COULD CALCULATE THE
UNION OF 2 QUERIES

WE COULD USE ONE QUERY INSIDE
ANOTHER (VIA SUBQUERIES)

BUT IN REALITY, QUERIES ARE PRETTY
PLUG-AND-PLAY
WE COULD CALCULATE THE
UNION OF 2 QUERIES

WE COULD USE ONE
QUERY INSIDE ANOTHER
(VIA SUBQUERIES)

WE COULD USE A
SUBQUERY TO POPULATE
A TABLE VIA INSERT

BUT IN REALITY, QUERIES ARE PRETTY PLUG-AND-PLAY

WE COULD USE ONE
QUERY INSIDE
ANOTHER (VIA
SUBQUERIES)

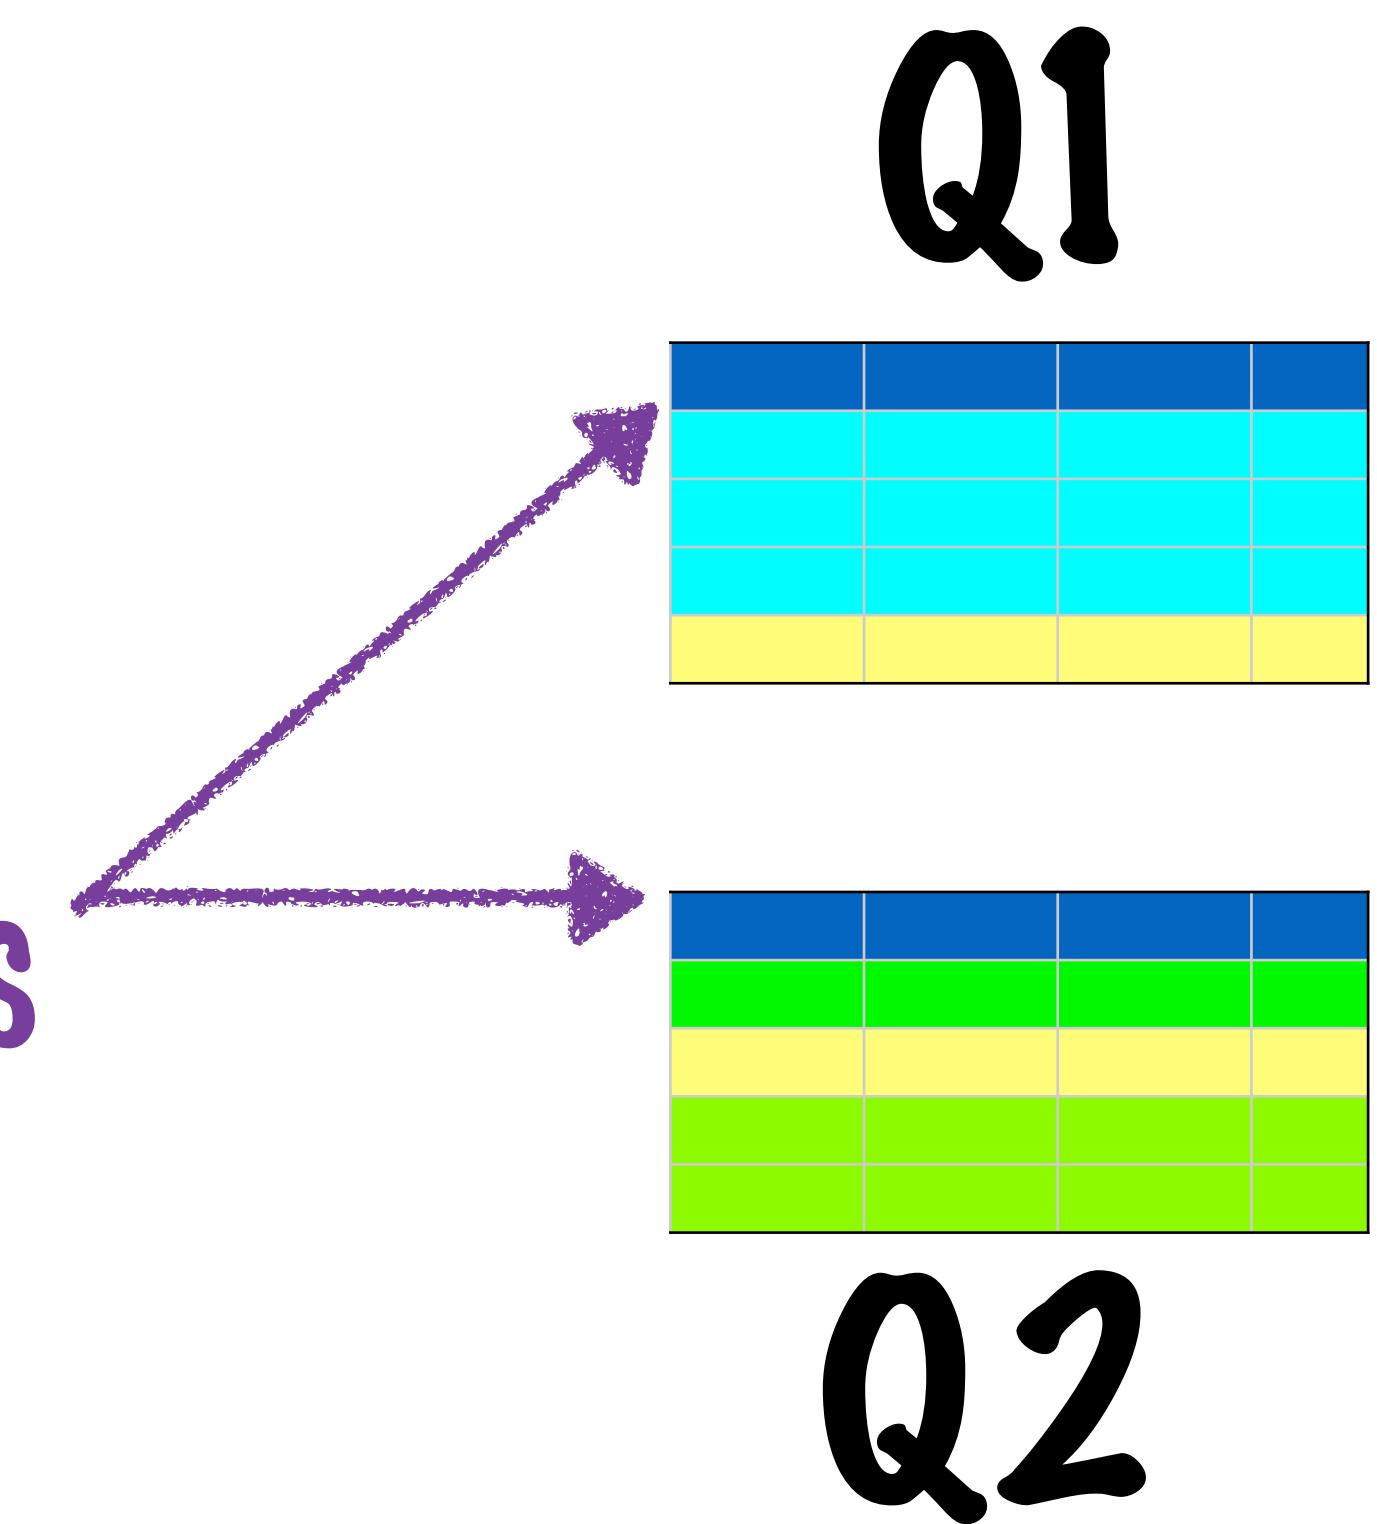
WE COULD
CALCULATE THE
UNION OF 2
QUERIES

WE COULD USE A
SUBQUERY TO
POPULATE A
TABLE VIA INSERT

WE COULD CALCULATE
THE UNION OF 2 QUERIES

PROVIDED THEY HAVE THE
SAME COLUMNS (NUMBER,
ORDER AND TYPE)

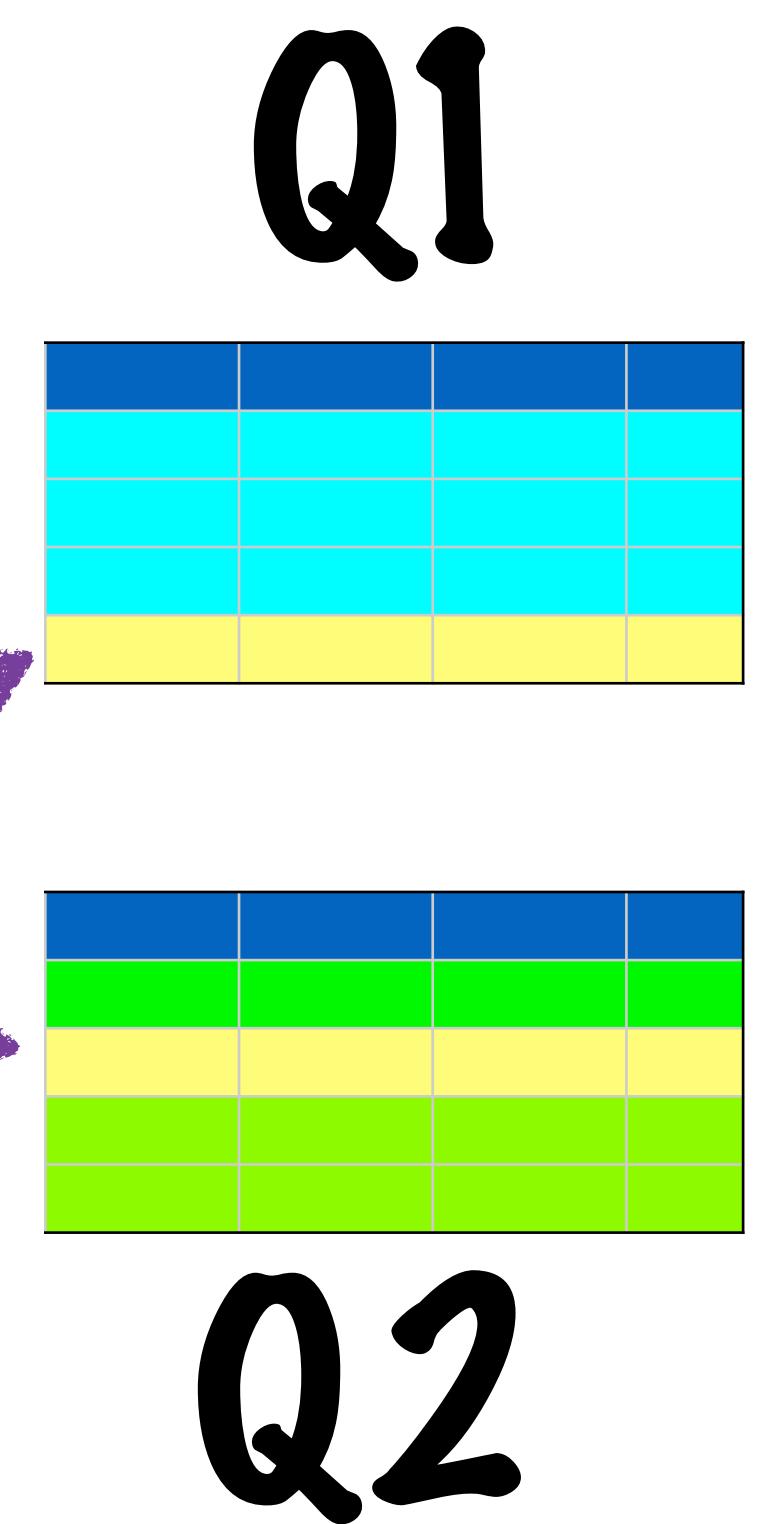
SAME
COLUMNS



WE COULD CALCULATE
THE UNION OF 2 QUERIES

PROVIDED THEY HAVE THE
SAME COLUMNS (NUMBER,
ORDER AND TYPE)

SOME COMMON
ROWS



WE COULD CALCULATE THE UNION OF 2 QUERIES

PROVIDED THEY HAVE THE SAME
COLUMNS (NUMBER, ORDER AND TYPE)

A QUERY IS A COMMAND
THAT RETURNS A TABLE
(ROWS AND COLUMNS)

Q1

Q2

WE COULD CALCULATE THE UNION OF 2 QUERIES

PROVIDED THEY HAVE THE SAME
COLUMNS (NUMBER, ORDER AND TYPE)

A QUERY IS A COMMAND THAT RETURNS
A TABLE (ROWS AND COLUMNS)

Q1

UNION



Q2

Q3

ONLY 1 COPY OF THE COMMON
ROWS (I.E. NO DUPLICATES)

WE COULD CALCULATE THE UNION OF 2 QUERIES

PROVIDED THEY HAVE THE SAME
COLUMNS (NUMBER, ORDER AND TYPE)

A QUERY IS A COMMAND THAT RETURNS
A TABLE (ROWS AND COLUMNS)

Q1

UNION ALL



Q2

Q4

UNION ALL MAINTAINS
DUPLICATES

WE COULD CALCULATE
THE UNION OF 2 QUERIES

ALSO THE INDIVIDUAL
QUERIES IN A UNION
CANNOT USE ORDER BY

WE COULD CALCULATE
THE UNION OF 2 QUERIES

HIVE DOES NOT SUPPORT

INTERSECT

MINUS

WHICH ARE SUPPORTED IN SQL

UNION

PET OWNERS

APT OWNERS

**SAME COLUMNS
(NUMBER, ORDER
AND TYPE - NAMES
COULD DIFFER)**

AptNumber	Name
123	John
345	Tim
349	Nikhil
567	Bilal

FlatNumber	Name
234	Mary
567	Bilal
897	Alan
903	Ellen

UNION

PET OWNERS

AptNumber	Name
123	John
345	Tim
349	Nikhil
567	Bilal

1 COMMON ROW

APT OWNERS

FlatNumber	Name
234	Mary
567	Bilal
897	Alan
903	Ellen

UNION

```
SELECT APTNUMBER, NAME  
FROM PetOwners
```

UNION

```
SELECT FLATNUMBER AS  
APTPNUMBER, NAME FROM  
AptOwners;
```

PET OWNERS

AptNumber	Name
123	John
345	Tim
349	Nikhil
567	Bilal

APT OWNERS

FlatNumber	Name
234	Mary
567	Bilal
897	Alan
903	Ellen

UNION

```
SELECT APTNUMBER, NAME  
FROM PetOwners
```

UNION

```
SELECT FLATNUMBER AS  
APTPNUMBER, NAME FROM  
AptOwners;
```

ONLY 1 COPY OF THE COMMON
ROWS (I.E. NO DUPLICATES)



AptNumber	Name
123	John
345	Tim
349	Nikhil
234	Mary
567	Bilal
897	Alan
903	Ellen

UNION

```
SELECT APTNUMBER, NAME  
FROM PetOwners
```

UNION ALL

```
SELECT FLATNUMBER AS  
APTPNUMBER, NAME FROM  
AptOwners;
```

PET OWNERS

AptNumber	Name
123	John
345	Tim
349	Nikhil
567	Bilal

APT OWNERS

FlatNumber	Name
234	Mary
567	Bilal
897	Alan
903	Ellen

UNION

```
SELECT APTNUMBER, NAME  
FROM PetOwners
```

UNION ALL

```
SELECT FLATNUMBER AS  
APTNUMBER, NAME FROM  
AptOwners;
```



UNION ALL MAINTAINS
DUPLICATES

AptNumber	Name
123	John
345	Tim
349	Nikhil
567	Bilal
234	Mary
567	Bilal
897	Alan
903	Ellen

UNION

(SELECT APTNUMBER, NAME FROM
PetOwners ORDER BY NAME)

WON'T WORK!!

(SELECT FLATNUMBER AS
APTPNUMBER, NAME FROM
AptOwners ORDER BY NAME);

PET
OWNERS

AptNumber	Name
123	John
345	Tim
349	Nikhil
567	Bilal

APT
OWNERS

FlatNumber	Name
234	Mary
567	Bilal
897	Alan
903	Ellen

UNION

THIS IS FINE!

SELECT APTNUMBER,

NAME FROM

PetOwners

UNION

SELECT FLATNUMBER AS

APTNUMBER, NAME FROM

AptOwners

: Order by APTNUMBER; ::

PET OWNERS

AptNumber	Name
123	John
345	Tim
349	Nikhil
567	Bilal

APT OWNERS

FlatNumber	Name
234	Mary
567	Bilal
897	Alan
903	Ellen

BUT IN REALITY, QUERIES ARE PRETTY PLUG-AND-PLAY

WE COULD USE ONE
QUERY INSIDE
ANOTHER (VIA
SUBQUERIES)

WE COULD
CALCULATE THE
UNION OF 2
QUERIES

WE COULD USE A
SUBQUERY TO
POPULATE A
TABLE VIA INSERT

BUT IN REALITY, QUERIES ARE PRETTY PLUG-AND-PLAY

.....
WE COULD USE ONE
QUERY INSIDE
ANOTHER (VIA
SUBQUERIES)
.....

 WE COULD
CALCULATE THE
UNION OF 2
QUERIES

WE COULD USE A
SUBQUERY TO
POPULATE A
TABLE VIA INSERT

WE COULD USE ONE
QUERY INSIDE ANOTHER
(VIA SUBQUERIES)

HIVE OFFERS LIMITED
SUPPORT FOR SUBQUERIES

WE COULD USE ONE
QUERY INSIDE ANOTHER
(VIA SUBQUERIES)

HIVE SUPPORTS SUBQUERIES

ONLY IN THE FROM CLAUSE

&

IN THE WHERE CLAUSE

WE COULD USE ONE QUERY INSIDE
ANOTHER (VIA SUBQUERIES)

HIVE SUPPORTS SUBQUERIES

IN THE WHERE CLAUSE

USING

IN/EXISTS

WE COULD USE ONE QUERY INSIDE
ANOTHER (VIA SUBQUERIES)

HIVE SUPPORTS SUBQUERIES

IN THE WHERE CLAUSE

THROUGH

IN/EXISTS

AND NOT THROUGH

“=” EQUALITY

IN/EXISTS

HIVE SUPPORTS SUBQUERIES

IN THE WHERE CLAUSE

THROUGH

IN

EXISTS

IN
‘WHERE IN’ RETURNS ROWS THAT
MATCHES VALUES IN A SUBQUERY

LET US SAY WE HAVE TWO TABLES

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

IN

'WHERE IN' RETURNS ROWS THAT MATCHES VALUES IN A SUBQUERY

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

```
SELECT NAMES . SYMBOL  
FROM NAMES  
WHERE NAMES . SYMBOL IN  
(SELECT SYMBOL FROM REVENUE) ;
```

IN

'WHERE IN' RETURNS ROWS THAT MATCHES VALUES IN A SUBQUERY

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

```
SELECT NAMES . SYMBOL  
FROM NAMES  
WHERE NAMES . SYMBOL IN  
(SELECT SYMBOL FROM REVENUE) ;
```

IN

'WHERE IN' RETURNS ROWS THAT MATCHES VALUES IN A SUBQUERY

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

```
SELECT NAMES . SYMBOL  
FROM NAMES  
WHERE NAMES . SYMBOL IN  
(SELECT SYMBOL FROM REVENUE) ;
```



SYMBOL
RIL
NESTLEIND

NOT IN

'WHERE IN' RETURNS ROWS THAT DOES NOT MATCH VALUES IN A SUBQUERY

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

```
SELECT NAMES . SYMBOL  
FROM NAMES  
WHERE NAMES . SYMBOL NOT IN  
(SELECT SYMBOL FROM REVENUE) ;
```

SYMBOL
TATA

NOT IN & IN

**IN/NOT IN SUBQUERIES ONLY SELECT
A SINGLE COLUMN**

IN IN/NOT IN SUBQUERIES ONLY SELECT A SINGLE COLUMN

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES



```
SELECT NAMES . SYMBOL
FROM NAMES
WHERE NAMES . SYMBOL IN
(SELECT SYMBOL FROM REVENUE) ;
```

SYMBOL
RIL
NESTLEIND

NAMES

IN/EXISTS

IN/EXISTS SUBQUERY ARE SUPPORTED IN
THE WHERE CLAUSE

IN

EXISTS

EXISTS

THE SQL EXISTS CONDITION IS USED IN COMBINATION WITH A SUBQUERY

```
SELECT *
FROM TABLE_NAME
WHERE EXISTS
SUBQUERY;
```

IF THE SUBQUERY RETURNS AT LEAST ONE ROW, THE CONDITION IS SAID TO BE MET

EXISTS

THE SQL EXISTS CONDITION IS USED IN COMBINATION WITH A SUBQUERY

IF THE SUBQUERY RETURNS AT LEAST ONE ROW, THE CONDITION IS SAID TO BE MET

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

EXISTS
**IF THE SUBQUERY RETURNS AT LEAST ONE
ROW, THE CONDITION IS SAID TO BE MET**

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

```
SELECT NAMES . SYMBOL
FROM NAMES
WHERE EXISTS
(SELECT TOTAL _ REVENUE FROM REVENUE
NAMES . name = REVENUE . name) ;
```

EXISTS

IF THE SUBQUERY RETURNS AT LEAST ONE ROW, THE CONDITION IS SAID TO BE MET

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

```
SELECT NAMES . SYMBOL  
FROM NAMES  
WHERE EXISTS  
(SELECT TOTAL _ REVENUE FROM REVENUE  
NAMES . name = REVENUE . name) ;
```

EXISTS

IF THE SUBQUERY RETURNS AT LEAST ONE ROW, THE CONDITION IS SAID TO BE MET

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

```
SELECT NAMES . SYMBOL  
FROM NAMES  
WHERE EXISTS  
(SELECT TOTAL _ REVENUE FROM REVENUE  
NAMES . name = REVENUE . name) ;
```

EXISTS

THERE IS ONE MORE CONDITION

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

THE SUBQUERY NEEDS TO BE CORRELATED

(SELECT TOTAL_REVENUE FROM REVENUE
NAMES.name = REVENUE.name);

THE SUBQUERY NEEDS TO BE CORRELATED

OUTER
QUERY

```
SELECT *
FROM TABLE_NAME
WHERE EXISTS
SUBQUERY;
```

A SUBQUERY IS
CORRELATED
WHEN IT USES A
VALUE FROM THE
OUTER QUERY

EXISTS
THE REFERENCED COLUMNS .name
SHOULD BE IN SMALL CAPS

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

THE SUBQUERY NEEDS TO BE CORRELATED

FROM NAMES
WHERE EXISTS
(SELECT TOTAL_REVENUE FROM REVENUE
NAMES.name = REVENUE.name);

THE SUBQUERY NEEDS TO BE CORRELATED

THE REFERENCED COLUMNS .name
SHOULD BE IN SMALL CAPS

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

IF YOU DON'T USE SMALL CAPS, YOU
WILL SEE THIS ERROR

(SELECT TOTAL REVENUE FROM REVENUE

FAILED: SemanticException [Error 10250]: Line 1:84 Invalid SubQuery expression 'SYMBOL': For Exists/Not Exists operator
SubQuery must be Correlated.

EXISTS
**IF THE SUBQUERY RETURNS AT LEAST ONE
ROW, THE CONDITION IS SAID TO BE MET**

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

```
SELECT NAMES . SYMBOL
FROM NAMES
WHERE EXISTS
(SELECT TOTAL _ REVENUE FROM REVENUE
NAMES . name = REVENUE . name) ;
```



SYMBOL
RIL
NESTLEIND

NOT EXISTS
**IF THE SUBQUERY DOES NOT RETURN AT LEAST
ONE ROW, THE CONDITION IS SAID TO BE MET**

TOTAL REVENUE	NAME	SYMBOL
\$ 4 Billion	Reliance	RIL
\$ 0.5 Billion	Nestle	NESTLEIND

REVENUE

SYMBOL	NAME	SECTOR
RIL	Reliance	ENERGY
NESTLEIND	Nestle	FMCG
TATA	TATA	AUTOMOBILE

NAMES

```
SELECT NAMES . SYMBOL
FROM NAMES
WHERE NOT EXISTS
(SELECT TOTAL _ REVENUE FROM REVENUE
NAMES . name = REVENUE . name) ;
```

SYMBOL
TATA

WE COULD USE ONE
QUERY INSIDE ANOTHER
(VIA SUBQUERIES)

LET'S SEE AN EXAMPLE

WE COULD USE ONE QUERY INSIDE ANOTHER (VIA SUBQUERIES)

'STORES'

StoreID	StoreLocation	City
1	Bellandur	Bangalore
2	Koramangala	Bangalore

'PRODUCTS'

ProductID	ProductName
1	Bananas
2	Milk
3	Nutella
4	Peanut Butter

'SALES_DATA'

StoreID	ProductID	Date	Revenue
1	1	January 18,2016	8,236.33
1	3	January 18,2016	7,455.67
1	4	January 18,2016	5,316.89
1	2	January 18,2016	2,433.76
2	1	January 18,2016	9,456.01
2	3	January 18,2016	3,644.33
2	4	January 18,2016	8,988.64

WE COULD USE ONE
QUERY INSIDE ANOTHER
(VIA SUBQUERIES)

IF SUCH A DATABASE
ACTUALLY EXISTED

WE COULD USE ONE
QUERY INSIDE ANOTHER
(VIA SUBQUERIES)

IF SUCH A DATABASE ACTUALLY EXISTED

YOU CAN BET THAT SOMEONE
WOULD PULL A REPORT FOR
ANNUAL REVENUE

WE COULD USE ONE
QUERY INSIDE ANOTHER
(VIA SUBQUERIES)

IF SUCH A DATABASE ACTUALLY EXISTED

YOU CAN BET THAT SOMEONE WOULD
PULL A REPORT FOR ANNUAL REVENUE

IN FACT, THERE WOULD BE A
REALLY BORING WEEKLY MEETING
AROUND THIS REPORT

'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

```
SELECT
    p.ProductName, YEAR(orderdate), SUM(revenue)
FROM
    Sales_Data s
INNER JOIN
    Products p
ON
    s.ProductID = p.ProductID
WHERE
    (p.ProductName = 'Peanut Butter' or p.ProductName =
    'Nutella') AND (YEAR(OrderDate) = 2016)
GROUP BY
    p.ProductName, YEAR(orderdate);
```

'SALES_DATA'

StoreID	ProductID	Date	Revenue

'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

'SALES_DATA'

StoreID	ProductID	Date	Revenue

```
SELECT
    p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
    Sales_Data s
INNER JOIN
    Products p
ON
    s.ProductID = p.ProductID
WHERE
    (p.ProductName = 'Peanut Butter' or p.ProductName =
    'Nutella') AND (YEAR(OrderDate) = 2016)
GROUP BY
    p.ProductName , YEAR(orderdate);
```

'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

```
SELECT
    p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
    Sales_Data s
INNER JOIN
    Products p
ON
    s.ProductID = p.ProductID
WHERE
    (p.ProductName = 'Peanut Butter' or p.ProductName =
    'Nutella') AND (YEAR(OrderDate) = 2016)
GROUP BY
    p.ProductName , YEAR(orderdate) ;
```

'SALES_DATA'

StoreID	ProductID	Date	Revenue

'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

```
SELECT
    p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
    Sales_Data s
INNER JOIN
    Products p
ON
    s.ProductID = p.ProductID
WHERE
    (p.ProductName = 'Peanut Butter' or p.ProductName =
    'Nutella') AND (YEAR(OrderDate) = 2016)
GROUP BY
    p.ProductName , YEAR(orderdate);
```

'SALES_DATA'

StoreID	ProductID	Date	Revenue

'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

```
SELECT
    p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
    Sales_Data s
INNER JOIN
    Products p
ON
    s.ProductID = p.ProductID
WHERE
    (p.ProductName = 'Peanut Butter' or p.ProductName =
    'Nutella') AND (YEAR(OrderDate) = 2016)
GROUP BY
    p.ProductName , YEAR(orderdate);
```

'SALES_DATA'

StoreID	ProductID	Date	Revenue

'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

```
SELECT
    p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
    Sales_Data s
INNER JOIN
    Products p
ON
    s
```

'SALES_DATA'

StoreID	ProductID	Date	Revenue

ERRM..THIS QUERY WILL NEED TO
BE RE-WRITTEN EACH YEAR :-|

'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

```
SELECT
    p.ProductName, YEAR(orderdate), SUM(revenue)
FROM
    Sales_Data s
INNER JOIN
    Products p
ON
    s.ProductID = p.ProductID
```

'SALES_DATA'

StoreID	ProductID	Date	Revenue

NOT IF WE USE A SUB-QUERY!

la')

0

'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

```
SELECT
    p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
    Sales_Data AS s
INNER JOIN
    Products AS p
ON
    s.ProductID = p.ProductID
WHERE
    (p.ProductName = 'Leanne Butcher' OR p.ProductName = 'Kwesi Lee'),
    AND
    (YEAR(s.orderdate) IN (SELECT max(YEAR(orderdate)) FROM Sales_Data ))
GROUP BY
    p.ProductName , YEAR(orderdate);
```

'DATA'

Revenue

WE JUST PLUGGED ONE QUERY
INTO ANOTHER!

'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

SELECT

p. Pr

FROM

Sale

INNE

Prod

ON

s. Pr

WHE

HERE “=” WILL NOT WORK!

(p.ProductName = 'Peanut Butter' or p.ProductName = 'Nutella')
AND

(YEAR(s.orderdate) = (SELECT max(YEAR(orderdate)) FROM Sales_Data))

GROUP BY

p.ProductName, YEAR(orderdate);

'DATA'

Revenue



'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

SELECT

p.ProductName , YEAR(orderdate) , SUM(revenue)

FROM

Sales

INNER

Produc

ON

s.Prod

WHERE

(p.Pro

AND

(YEAR(s.orderdate) in (SELECT max(YEAR(orderdate)) FROM Sales_Data))

GROUP BY

p.ProductName , YEAR(orderdate) ;

'DATA'

Revenue

HERE, THE INNER QUERY RETURNS
A SINGLE VALUE..

'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT SOMEONE WOULD PULL A REPORT FOR ANNUAL REVENUE

```
SELECT
    p.ProductName, YEAR(orderdate), SUM(revenue)
FROM
    Sales
INNER JOIN
    Products p
ON
    s.ProductID = p.ProductID
WHERE
    (p.ProductName = 'Peanut Butter' OR p.ProductName = 'Nutella')
    AND
    (YEAR(s.orderdate)) IN (SELECT MAX(YEAR(orderdate)) FROM Sales_Data)
GROUP BY
    p.ProductName, YEAR(orderdate);
```

'DATA'

Revenue

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

```
SELECT
    p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
    Sales_Data s
INNER JOIN
    Products p
ON
    s.ProductID = p.ProductID
WHERE
    (p.ProductName = 'Peanut Butter' or p.ProductName = 'Nutella')
AND
    (YEAR(s.orderdate) in (SELECT max(YEAR(orderdate)) FROM Sales_Data))
GROUP BY
    p.ProductName , YEAR(orderdate);
```

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

ProductID	ProductName

'SALES_DATA'

StoreID	ProductID	Date	Revenue

LET'S SAY WE DON'T
WANT TO HARD CODE
THE PRODUCTS,
INSTEAD WE HAVE A
TOP SELLERS TABLE

```
s.ProductID = p.ProductID  
WHERE  
(p.ProductName = 'Peanut Butter' or p.ProductName = 'Nutella')  
AND  
(YEAR(s.orderdate) in (SELECT max(YEAR(orderdate)) FROM  
Sales_Data ))  
GROUP BY  
p.ProductName, YEAR(orderdate);
```

YOU CAN BET THAT
WE'D PULL A
ANNUAL
REPORT
BY QUARTER
(Revenue)

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

ProductID	ProductName

'SALES_DATA'

StoreID	ProductID	Date	Revenue

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

productID	ProductName

WE MIGHT WANT TO DO
SOMETHING LIKE THIS

'SALES_DATA'

Date	Revenue

WHERE

(p.ProductName IN (SELECT ProductName FROM Top_Sellers))

AND

(YEAR(s.orderdate) in (SELECT max(YEAR(orderdate)) FROM Sales_Data))

GROUP BY

p.ProductName, YEAR(orderdate);

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

ProductID	ProductName

'SALES DATA'

orderdate	Revenue

SELECT

UNFORTUNATELY, THIS IS NOT
SUPPORTED IN HIVE

WHERE

```
(p.ProductName IN (SELECT ProductName FROM Top_Sellers))  
AND  
(YEAR(s.orderdate) in (SELECT max(YEAR(orderdate)) FROM  
Sales_Data ))  
GROUP BY  
p.ProductName, YEAR(orderdate);
```

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

ProductID	ProductName

S_DATA

Revenue

THIS IS THE ERROR MESSAGE IN
HIVE 2.0.0

```
SELECT
  p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
Sales
INNER
Prod
ON
S.P
WHERE
(p.ProductName IN (SELECT ProductName FROM Top_Sellers))
AND
(YEAR(s.orderdate) in (SELECT max(YEAR(orderdate)) FROM
Sales Data ))
FAILED: SemanticException [Error 10249]: Line 12:19 Unsupported SubQuery Expression 'orderdate': Only 1 SubQuery expression is supported.
p.ProductName , YEAR(orderdate);
```

'PRODUCTS'

ProductID	ProductName

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL

SELECT
p.ProductName , YEAR (orderdate) , SUM (revenue)

FROM

Sale

INNE

Pro

ON

s . P

WHERE

(p.ProductName IN (SELECT ProductName FROM Top_Sellers))

AND

(YEAR(s.orderdate) in (SELECT max(YEAR(orderdate)) FROM Sales_Data))

FAILED: SemanticException [Error 10249]: Line 12:19 Unsupported SubQuery Expression 'orderdate': Only 1 SubQuery expression is supported.

'TOP SELLERS'

ProductID	ProductName

-DATA-

Revenue

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

ProductID	ProductName

```
SELECT  
p.ProductName , YEAR(orderdate) , SUM(revenue)  
FROM  
Sales_Data s
```

IN THIS QUERY WE CAN FIX THIS IN A
SLIGHTLY DIFFERENT WAY, YOU'LL SEE THIS
LATER IN THIS CLASS

```
(YEAR(s.orderdate) in (SELECT max(YEAR(orderdate)) FROM
```

FAILED: SemanticException [Error 10249]: Line 12:19 Unsupported SubQuery Expression 'orderdate': Only 1 SubQuery expression is supported.

```
p.ProductName , YEAR(orderdate);
```

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

ProductID	ProductName

'SALES_DATA'

StoreID	ProductID	Date	Revenue

TIME PASSES, AND OUR SALES_DATA
TABLE BECOMES ENORMOUS - TOO
BIG TO PULL LIKE THIS

```
SELECT
    p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
    Sales_Data s
INNER JOIN
    Products p
ON
    s.ProductID = p.ProductID
WHERE
    (YEAR(orderdate) = 2015)
    AND (p.CategoryID = 1)
    AND (s.StoreID = 1)
GROUP BY
    p.ProductName , YEAR(orderdate)
```

**YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE**

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

ProductID	ProductName

```
SELECT  
p.ProductName , YEAR(orderdate) , SUM(revenue)  
FROM
```

```
(select * from Sales_Data where ProductID in (select  
PRODUCTID from Top Sellers)) s
```

```
INNER JOIN
```

```
Products p
```

```
ON
```

```
s.Pro
```

```
WHERE
```

```
(YEAR
```

```
Sales
```

```
GROU
```

```
p.Pro
```

'SALES_DATA'

**WE JUST PLUGGED ONE QUERY
INTO ANOTHER!**

Revenue

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

ProductID	ProductName

```
SELECT  
p.ProductName , YEAR(orderdate) , SUM(revenue)  
FROM
```

```
(select * from Sales_Data where ProductID in (select  
PRODUCTID from Top_Sellers)) s
```

```
INNER JOIN
```

```
Produ
```

```
ON
```

```
s.Pro
```

```
WHERE
```

```
(YEAR
```

```
Sales
```

```
GROUP
```

```
p.Productname , YEAR(oraerdate);
```

'SALES_DATA'

HERE, THE INNER QUERY IS A
SUBSET OF A TABLE...

Revenue

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

```
SELECT  
p.ProductName , YEAR(orderdate) , SUM(revenue)  
FROM
```

```
(select * from Sales_Data where ProductID in (select  
PRODUCTID from Top_Sellers)) s
```

```
INNER JOIN
```

```
Product
```

```
ON
```

```
s.Pro
```

```
WHERE
```

```
(YEAR
```

```
Sales
```

```
GROUP
```

```
p.Pro
```

'SALES_DATA'

THE OUTER QUERY USES THAT
TABLE SUBSET (ALSO A TABLE) IN
THE FROM CLAUSE

Revenue

Revenue

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

ProductID	ProductName

```
SELECT
    p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
    (select * from Sales_Data where ProductID in (select
    PRODUCTID from Top_Sellers)) s
INNER JOIN
    Products p
ON
    s.ProductID = p.ProductID
WHERE
    (YEAR(orderdate) = 2015)
    AND (p.CategoryID = 1)
GROUP BY
    p.ProductName
```

'SALES_DATA'

THE OUTER QUERY ASSESSES THE
RESULTS OF THE SUBQUERY LIKE IT
DOES WITH A TABLE

YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

ProductID	ProductName

```
SELECT
    p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
    (select * from Sales_Data where ProductID in (select PRODUCTID
from Top_Sellers)) s
INNER JOIN
    Products p
ON
    s.ProductID = p.ProductID
WHERE
    (YEAR(orderdate) = 2015)
Sales_Data GROUP BY
    p.ProductName , YEAR(orderdate)
```

'SALES_DATA'

THE SUB-QUERY QUERY MUST BE
GIVEN AN ALIAS. HERE 'S' IS THE
SUBQUERY ALIAS

Revenue

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

ProductID	ProductName

'SALES_DATA'

```
SELECT
    p.ProductName, YEAR(orderdate), SUM(revenue)
FROM
    (select * from Sales_Data where ProductID in (select PRODUCTID
from Top_Sellers)) s
INNER JOIN
    Product p
ON
    s.ProductID = p.ProductID
WHERE
    (YEAR(orderdate) - YEAR(salesdate)) >= 1
Sales_Data
GROUP BY
    p.ProductName, YEAR(orderdate),
```

THIS ALSO FIXES OUR OLDER ISSUE
WHERE WE COULD NOT USE MORE
THAN ONE SUBQUERY IN A TABLE

Revenue

Revenue

**YOU CAN BET THAT
SOMEONE WOULD PULL A
REPORT FOR ANNUAL
REVENUE**

'PRODUCTS'

ProductID	ProductName

'TOP SELLERS'

```
SELECT
    p.ProductName , YEAR(orderdate) , SUM(revenue)
FROM
    (select * from Sales_Data where ProductID in (select
    PRODUCTID from Top_Sellers)) s
```

**THE QUERY IS NOW ENTIRELY FREE OF
HARDCODED VALUES OR VERY LARGE
INTERMEDIATE TABLES - THANKS TO THE
USE OF SUBQUERIES!**

WE COULD USE ONE
QUERY INSIDE ANOTHER
(VIA SUBQUERIES)

IF SUCH A DATABASE ACTUALLY EXISTED

YOU CAN BET THAT SOMEONE WOULD
PULL A REPORT FOR ANNUAL REVENUE

IN FACT, THERE WOULD BE A REALLY BORING
WEEKLY MEETING AROUND THIS REPORT

THE QUERY IS NOW ENTIRELY FREE OF HARDCODED
VALUES OR VERY LARGE INTERMEDIATE TABLES -
THANKS TO THE USE OF SUBQUERIES!

WE COULD USE ONE
QUERY INSIDE ANOTHER
(VIA SUBQUERIES)

'STORES'

StoreID	StoreLocation	City
1	Bellandur	Bangalore
2	Koramangala	Bangalore

'PRODUCTS'

ProductID	ProductName
1	Bananas
2	Milk
3	Nutella
4	Peanut Butter

SUBQUERIES REALLY SIMPLIFY &
ROBUSTIFY QUERIES :-)

'SALES_DATA'

StoreID	ProductID	Date	Revenue
1	1	January 18,2016	8,236.33
1	3	January 18,2016	7,455.67
1	4	January 18,2016	5,316.80

WE COULD USE ONE
QUERY INSIDE ANOTHER
(VIA SUBQUERIES)

SUBQUERIES REALLY SIMPLIFY &
ROBUSTIFY QUERIES :-)

HIVE SUPPORTS SUBQUERY
IN THE WHERE/FROM CLAUSE
THROUGH
IN
EXISTS

BUT IN REALITY, QUERIES ARE PRETTY PLUG-AND-PLAY

.....
WE COULD USE ONE
QUERY INSIDE
ANOTHER (VIA
SUBQUERIES)
.....

 WE COULD
CALCULATE THE
UNION OF 2
QUERIES

WE COULD USE A
SUBQUERY TO
POPULATE A
TABLE VIA INSERT

BUT IN REALITY, QUERIES ARE PRETTY PLUG-AND-PLAY

- WE COULD USE ONE QUERY INSIDE ANOTHER (VIA SUBQUERIES)
- WE COULD CALCULATE THE UNION OF 2 QUERIES
- WE COULD USE A SUBQUERY TO POPULATE A TABLE VIA INSERT

WE COULD USE A SUBQUERY TO
POPULATE A TABLE VIA INSERT

WE HAVE ALREADY COME ACROSS THE
INSERT & CREATE TABLE STATEMENTS

RECAP

LET'S SAY WE WANT TO CREATE TABLE WHICH STORES LIST OF PRODUCTS ALONG WITH THE STORE LOCATIONS

CREATE A TABLE FIRST

```
CREATE TABLE stores_product_data  
(  
StoreLocation VARCHAR(30) ,  
Product VARCHAR(30)  
);
```

StoreLocation	Product
Bellandur	Bananas
Bellandur	Nutella
Bellandur	Peanut Butter
Bellandur	Milk
Koramangala	Bananas
Koramangala	Nutella
Koramangala	Peanut Butter
Koramangala	Milk

StoreLocation	Product	Date	Revenue

RECAP

INSERT INTO THAT TABLE

StoreLocation	Product
Bellandur	Bananas
Bellandur	Nutella
Bellandur	Peanut Butter
Bellandur	Milk
Koramangala	Bananas
Koramangala	Nutella
Koramangala	Peanut Butter
Koramangala	Milk

```
insert overwrite table stores_product_data
select StoreLocation, Product
from Sales_Data;
```

StoreLocation	Product	Date	Revenue

**WE USED A SUBQUERY TO
POPULATE A TABLE VIA INSERT**

**SUBQUERIES ARE FAR MORE
CONVENIENT TO POPULATE A TABLE
FROM EXISTING TABLES IN THE
DATABASE**

BUT SUBQUERIES ARE FAR MORE
CONVENIENT TO POPULATE A TABLE
FROM EXISTING TABLES IN THE DATABASE

WE COULD EITHER -

- CREATE THE TABLE AS USUAL, THEN
INSERT DATA USING A SUBQUERY
- CREATE THE TABLE AND POPULATE
DIRECTLY USING A SUBQUERY

- CREATE THE TABLE AS USUAL, THEN
INSERT DATA USING A SUBQUERY

```
CREATE TABLE EmailAddresses
(
Email VARCHAR(30),
Category VARCHAR(10)
);
```

```
INSERT INTO EmailAddresses
SELECT distinct Email, "Student" AS Category FROM
Students;
```

- CREATE THE TABLE AS USUAL, THEN
INSERT DATA USING A SUBQUERY

```
CREATE TABLE EmailAddresses  
(  
Email VARCHAR(30),  
Category VARCHAR(10)  
);
```

FIRST CREATE THE TABLE
EXACTLY AS USUAL

```
INSERT INTO EmailAddresses  
SELECT distinct Email, "Student" AS Category FROM  
Students;
```

- CREATE THE TABLE AS USUAL, THEN INSERT DATA USING A SUBQUERY

```
CREATE TABLE EmailAddresses  
(  
Email VARCHAR(30),  
Category VARCHAR(10)  
);
```

NEXT USE A DIFFERENT FORM
OF THE INSERT STATEMENT,
THAT IS FOLLOWED BY A QUERY

```
INSERT INTO EmailAddresses  
SELECT distinct Email, "Student" AS Category FROM  
Students;
```

- CREATE THE TABLE AS USUAL, THEN INSERT DATA USING A SUBQUERY

```
CREATE TABLE EmailAddresses  
(  
Email VARCHAR(30),  
Category VARCHAR(10)  
);
```

THIS INSERT STARTS OFF AS
USUAL, WITH THE NAME OF
THE TABLE TO INSERT INTO

```
INSERT INTO EmailAddresses  
SELECT distinct Email, "Student" AS Category FROM  
Students;
```

- CREATE THE TABLE AS USUAL, THEN
INSERT DATA USING A SUBQUERY

```
CREATE TABLE EmailAddresses
(
Email VARCHAR(30),
Category VARCHAR(10)
);
```

THIS INSERT STARTS OFF AS
USUAL, WITH THE NAME OF
THE TABLE TO INSERT INTO

```
INSERT INTO EmailAddresses
SELECT distinct Email, "Student" AS Category FROM
Students;
```

- CREATE THE TABLE AS USUAL, THEN
INSERT DATA USING A SUBQUERY

```
CREATE TABLE EmailAddresses  
(  
Email VARCHAR(30),  
Category VARCHAR(10)  
);
```

BUT WHAT FOLLOWS
IS A QUERY

```
INSERT INTO EmailAddresses  
SELECT distinct Email, "Student" AS Category FROM  
Students;
```

- CREATE THE TABLE AS USUAL, THEN
INSERT DATA USING A SUBQUERY

```
CREATE TABLE EmailAddresses  
(  
Email VARCHAR(30),  
Category VARCHAR(10)  
);
```

BUT WHAT FOLLOWS
IS A QUERY

ALL THAT'S NEEDED IS FOR THE
QUERY TO MATCH THE TABLE IN THE
NUMBER AND TYPE OF COLUMNS

```
INSERT INTO EmailAddresses  
SELECT distinct Email, "Student" AS Category FROM  
Students;
```

BUT SUBQUERIES ARE FAR MORE
CONVENIENT TO POPULATE A TABLE
FROM EXISTING TABLES IN THE DATABASE

WE COULD EITHER -

- CREATE THE TABLE AS USUAL, THEN
INSERT DATA USING A SUBQUERY
- CREATE THE TABLE AND POPULATE
DIRECTLY USING A SUBQUERY

- CREATE THE TABLE AND POPULATE DIRECTLY USING A SUBQUERY

```
CREATE TABLE EmailAddresses
(
Email VARCHAR(30),
Category VARCHAR(10)
)
AS
SELECT distinct Email, "Student" AS
Category FROM Students;
```

- CREATE THE TABLE AND POPULATE DIRECTLY USING A SUBQUERY

```
CREATE TABLE EmailAddresses  
(  
Email VARCHAR(30),  
Category VARCHAR(10)  
)  
AS
```

```
SELECT distinct Email, "Student" AS Category  
FROM Students;
```

TABLE
DEFINITION

QUERY

- CREATE THE TABLE AND POPULATE DIRECTLY USING A SUBQUERY

```
CREATE TABLE EmailAddresses TABLE
```

```
(  
Email  
Category  
)
```

```
AS
```

```
SELECT distinct Email, "Student" AS  
Category FROM Students;
```

UNFORTUNATELY, THIS DOESN'T
WORK IN HIVE-2.0.0

- CREATE THE TABLE AND POPULATE
DIRECTLY USING A SUBQUERY

**WHEN YOU RUN THIS QUERY, YOU
GET THIS RESULT**

CREATE

(

Email

Category VARCHAR(10)

)

LINKED BY 'AS'

QUERY

```
hive> CREATE TABLE EmailAddresses
> (
> Email VARCHAR(30),
> Category VARCHAR(10)
> )
> AS
> SELECT distinct Email,"Student" AS Category FROM Students;
```

```
FAILED: SemanticException [Error 10065]: CREATE TABLE AS SELECT command cannot specify the list of columns for the target table
```

THIS COMMAND IS CALLED
“CREATE-TABLE-AS-SELECT”(CTAS)
COMMAND

CREATE
(
Email
Category VARCHAR(10)
)

LINKED BY ‘AS’

```
hive> CREATE TABLE EmailAddresses
> (
> Email VARCHAR(30),
> Category VARCHAR(10)
> )
> AS
> SELECT distinct Email,"Student" AS Category FROM Students;
```

FAILED: SemanticException [Error 10065]: CREATE TABLE AS SELECT command cannot specify the list of columns for the target table

FROM HIVE DOCUMENTATION

Create Table As Select (CTAS)

Tables can also be created and populated by the results of a query in one create-table-as-select (CTAS) statement. The table created by CTAS is atomic, meaning that the table is not seen by other users until all the query results are populated. So other users will either see the table with the complete results of the query or will not see the table at all.

There are two parts in CTAS, the SELECT part can be any [SELECT statement](#) supported by HiveQL. The CREATE part of the CTAS takes the resulting schema from the SELECT part and creates the target table with other table properties such as the SerDe and storage format.

CTAS has these restrictions:

- The target table cannot be a partitioned table.
- The target table cannot be an external table.
- The target table cannot be a list bucketing table.

Example:

```
CREATE TABLE new_key_value_store
  ROW FORMAT SERDE "org.apache.hadoop.hive.serde2.columnar.ColumnarSerDe"
  STORED AS RCFILE
  AS
  SELECT
    FROM key
  SORT BY
```

LET US SEE SOMETHING WHICH ACTUALLY WORKS

- CREATE THE TABLE AND POPULATE DIRECTLY USING A SUBQUERY

TABLE
DEFINITION

```
CREATE TABLE EmailAddresses  
AS  
SELECT distinct  
Email, "Student" AS Category  
FROM Students;
```

- CREATE THE TABLE AND POPULATE DIRECTLY USING A SUBQUERY

CREATE TABLE EmailAddresses TABLE
DEFINITION
AS

SELECT distinct Email 'Student' AS
Cat
UNI
SEL
Cat

**SKIP THE COLUMN DEFINITIONS
AND IT WILL WORK**

BUT SUBQUERIES ARE FAR MORE
CONVENIENT TO POPULATE A TABLE
FROM EXISTING TABLES IN THE DATABASE

WE COULD EITHER -

- CREATE THE TABLE AS USUAL, THEN
INSERT DATA USING A SUBQUERY
- CREATE THE TABLE AND POPULATE
DIRECTLY USING A SUBQUERY

**SUBQUERIES ARE FAR MORE
CONVENIENT TO POPULATE A
TABLE FROM EXISTING TABLES IN
THE DATABASE**

BUT IN REALITY, QUERIES ARE PRETTY PLUG-AND-PLAY

- WE COULD USE ONE QUERY INSIDE ANOTHER (VIA SUBQUERIES)
- WE COULD CALCULATE THE UNION OF 2 QUERIES
- WE COULD USE A SUBQUERY TO POPULATE A TABLE VIA INSERT

BUT IN REALITY, QUERIES ARE PRETTY PLUG-AND-PLAY

- WE COULD USE ONE QUERY INSIDE ANOTHER (VIA SUBQUERIES)
- WE COULD CALCULATE THE UNION OF 2 QUERIES
- WE COULD USE A SUBQUERY TO POPULATE A TABLE VIA INSERT